

MRV-EffScreen: Temporal Generalization and Consistency Screening of Ship Energy-Efficiency Labels from Public THETIS-MRV Emission Reports

Bo Li ¹ 

¹ China Maritime Service Center, Beijing, China; li.bo@cmaritime.com.cn

Abstract

Public ship-level emission reports provide a transparent basis for data-driven maritime decarbonization studies, but predictive analyses must avoid target leakage and must be evaluated under realistic temporal shifts. This study proposes MRV-EffScreen, a reproducible benchmark framework for relative ship energy-efficiency stratification and manual emission-consistency review using public EU THETIS-MRV emission reports. We harmonized 2018–2024 public records, constructed ship-type-year tertile labels from CO₂ emissions per distance, and evaluated non-leakage static and non-emission operational features under a temporal design. The processed dataset contained 90,745 records, 22,551 unique IMO numbers, and 17 ship types; the main 2018–2023 experiment contained 75,580 records, of which 73,045 were labeled after applying a minimum 30-row ship-type-year labeling rule. Fixed-parameter baseline models were trained on 2018–2021 records, reported on a 2022 temporal holdout, tested on 2023 records, and checked on 2024 full-year reports. Histogram-based gradient boosting with non-emission operational features achieved Macro-F1 = 0.603 and balanced accuracy = 0.605 on the 2023 temporal test set, while the best random stratified split reached Macro-F1 = 0.657. An IMO-overlap audit showed that 82.8% of 2023 test IMO numbers also appeared in the training years, so the temporal result should be read as year-ahead repeated-fleet generalization rather than new-vessel generalization. A separate consistency-review module combining Isolation Forest, Local Outlier Factor, and regression residual ranking prioritized 226 all-method review candidates at a 2% threshold. Public candidate tables remove direct vessel identifiers and are intended for manual inspection, not compliance adjudication.

Keywords: green shipping; THETIS-MRV; ship energy efficiency; temporal generalization; machine learning; consistency screening; emission consistency; maritime decarbonization

1. Introduction

International shipping is central to global trade and remains a difficult-to-decarbonize transport sector. The Fourth IMO Greenhouse Gas Study reported that greenhouse gas emissions from shipping increased between 2012 and 2018 and that shipping accounted for a measurable share of global anthropogenic emissions [5]. The 2023 IMO greenhouse gas strategy further strengthened the long-term ambition for international shipping to reach net-zero greenhouse gas emissions by or around 2050, with indicative checkpoints for 2030 and 2040 [6]. Operational instruments such as the Energy Efficiency Existing Ship Index (EEXI) and Carbon Intensity Indicator (CII) also reflect a broader regulatory shift from isolated technical standards toward repeated annual monitoring and operational carbon-intensity assessment [7].

The European Union monitoring, reporting, and verification (EU MRV) system is one of the most important public ship-level emission data sources. THETIS-MRV publishes annual emission reports for ships covered by the MRV regulation, including CO₂ emissions, fuel consumption, distance travelled, time spent at sea, cargo-related variables, ship type, and technical efficiency information [1–3]. Regulation (EU) 2023/957 further links maritime MRV to the EU Emissions Trading System

and extends the monitoring context to additional greenhouse gases and ship types [4]. These records are not voyage-level operational traces, but they provide official, structured, and annually repeated ship-level observations. This makes the dataset particularly valuable for reproducible studies on maritime emissions and green-shipping decision support.

Previous work has used EU MRV data to describe ship emissions, assess regulatory implications, and review early patterns after the mechanism's implementation. Bullock et al. showed how the early release of EU MRV data enabled ship-level analysis of committed shipping emissions [11]. Yan et al. developed quantitative MRV-based fuel-consumption prediction models using early MRV records and external ship features [8]. Luo et al. reviewed five years of MRV mechanism application and argued that the public data remain under-exploited for parameter analysis, temporal interpretation, and future policy support [9]. Xing et al. recently used MRV data to review ship energy-efficiency framework issues and operational-profile differences among ship types [10]. These studies establish the value of MRV data, but there remains a practical methodological gap: how can public MRV tables be used for predictive energy-efficiency stratification without leaking the target variable, and how should models be evaluated when deployment would require generalization to later reporting years?

This study addresses three research questions: whether non-leakage public MRV features can support moderate year-ahead relative CO₂-intensity stratification; how much random stratified evaluation overstates performance relative to temporal evaluation; and whether related emission, fuel, and time fields can support a cautious manual consistency-review queue without being interpreted as regulatory findings.

Target leakage is a central risk in this setting. If a label is derived from CO₂ emissions per distance, then total CO₂ emissions, fuel consumption, fuel per distance, and related intensity variables are formulaically or physically close to the target. Including them in the main classification features would produce overly optimistic results and would obscure whether the model learns any deployable non-leakage signal. Leakage has been recognized as a general source of invalid performance estimates in data mining and predictive modeling [17]. For public regulatory datasets, this issue is especially important because misleadingly strong results can lead to inappropriate interpretations in policy-facing contexts.

This study therefore separates two tasks. The first task is relative efficiency stratification using only non-leakage static and non-emission operational features. The second task is emission consistency screening, where fuel, CO₂, distance, time, and intensity fields are intentionally used to identify records that deserve manual inspection. The second task is not framed as fraud detection, violation detection, or compliance adjudication.

The contributions of this study are as follows:

1. A reproducible 2018–2024 THETIS-MRV processing workflow is built, with 2018–2023 used as the stable-schema main experiment and 2024 full-year reports retained as an external-year check.
2. A ship-type-year tertile label is defined from CO₂ emissions per distance to support relative efficiency stratification within comparable fleet-year groups.
3. Non-leakage feature sets are evaluated under temporal extrapolation and random stratified evaluation, showing that random splits overestimate generalization.
4. Unified and ship-type-specific models are compared for the five largest ship types, revealing selective but not universal gains from specialization.
5. A separate consistency-screening module combines Isolation Forest, Local Outlier Factor, and regression residual ranking to prioritize manual consistency-review candidates.

2. Related Work

2.1. Maritime Decarbonization and MRV Data

Maritime decarbonization depends on both technological change and reliable operational monitoring. IMO policy has increasingly emphasized vessel-level and annual carbon-intensity assessment, while EU MRV provides a public reporting infrastructure for ship-level emissions and related op-

erational parameters [1,2,4,6]. Compared with aggregate inventories, MRV records provide much finer ship-level granularity. Compared with commercial vessel databases or AIS-derived proprietary products, the public MRV releases are open and reproducible, although their annual aggregation limits the ability to model voyage-level drivers such as weather, speed profiles, port congestion, and loading conditions.

Recent MRV studies show a clear progression from descriptive analysis toward quantitative modeling. Yan et al. used MRV records and additional vessel attributes to predict annual average fuel consumption by ship type [8]. Luo et al. synthesized the first five years of MRV operation and emphasized both policy relevance and remaining data-use gaps [9]. Xing et al. further examined MRV-reported operational profiles and energy-efficiency indicators under the evolving technical and operational efficiency framework [10]. In parallel, machine-learning studies using onboard or in-service data have shown the value of tree ensembles and related models for ship energy-consumption prediction [12]. MRV-EffScreen differs from these lines by using only public annual MRV tables, by explicitly separating non-leakage classification from consistency screening, and by treating temporal generalization as the main evaluation target.

2.2. Relative Ship Efficiency Labels and Fleet Heterogeneity

Ship energy efficiency cannot be compared meaningfully without considering ship type, scale, and operating context. Bulk carriers, tankers, container ships, passenger ships, and specialized vessels differ in design, cargo function, and normal operating profiles. The proposed label is therefore constructed within each ship-type-year group rather than across all ships. This choice is not a replacement for EEXI, CII, or any regulatory score. Instead, it creates a reproducible relative task inside the public MRV data.

Fleet heterogeneity also affects modeling strategy. A unified model benefits from larger training data and can learn cross-ship-type regularities. Ship-type-specific models may capture category-specific patterns but use fewer records and may overfit. This trade-off motivates the ship-type ablation performed in this study.

2.3. Leakage Control and Consistency-Review Screening

The same MRV variables can play different roles depending on the task. CO₂ and fuel variables are unsuitable as features for predicting a label derived from CO₂ intensity, but they are relevant for consistency screening. This distinction is central to MRV-EffScreen. The classification feature sets exclude target-family variables; the consistency-screening feature set uses them only in an unsupervised or residual-screening context.

Isolation Forest isolates unusual observations through random partitioning and is well suited to large tabular datasets [15]. Local Outlier Factor identifies observations whose local density differs from that of neighboring points [16]. Regression residual ranking provides a complementary view by measuring whether a reported CO₂-per-distance value is difficult to reconstruct from related consistency fields. Combining these methods reduces dependence on a single screening score. Because no external audit labels are available, this branch is designed as a manual review queue rather than a validated anomaly detector.

3. Materials and Methods

3.1. Data Source and Scope

This study uses public emission reports released through EU THETIS-MRV. Raw workbooks were downloaded by reporting year and versioned locally with interface metadata and SHA256 checksums. The 2018–2023 workbooks share a stable 62-column schema. The 2024 workbook expands to 113 columns, adding company, methane, nitrous oxide, CO₂-equivalent, and EU ETS-related fields. To preserve comparability, the main predictive experiment uses the stable 2018–2023 schema, while 2024 full-year emission reports are used as an external-year check. Partial 2024 emission reports are

excluded from the main evaluation because their reporting periods do not represent complete annual operation.

Table 1. Dataset coverage and temporal split statistics.

Quantity	Value
Processed records, 2018–2024	90,745
Main-experiment records, 2018–2023	75,580
Main-experiment labeled records	73,045
Main-experiment unlabeled records	2,535
Labeling threshold	At least 30 records per ship-type-year group
Ship-type-year groups below threshold, 2018–2023	7
2024 full-year external records	14,139
2024 partial records excluded	1,026
Unique IMO numbers	22,551
Ship types	17
Train, 2018–2021 labeled records	47,380
2022 temporal holdout labeled records	13,050
2023 temporal test labeled records	12,615
External check, 2024 full-year labeled records	12,905




Figure 1. Annual coverage of THETIS-MRV public emission reports from 2018 to 2024.

3.2. Relative Efficiency Label

The main target is a three-class relative efficiency label derived from the CO₂-per-distance variable. Within each ship_type and reporting_year group, records are ranked by CO₂ emissions per distance and divided into tertiles: efficient, medium, and inefficient. Labels are assigned only when the corresponding ship-type-year group contains at least 30 records with non-missing CO₂ per distance. Groups below this threshold and rows with missing target values remain unlabeled and are excluded from the classification dataset. In the 2018–2023 main experiment, 73,045 of 75,580 records are labeled; the remaining 2,535 records comprise 69 records from below-threshold groups and 2,466 records with missing target values. This construction avoids direct comparison of raw CO₂ intensity across fundamentally different ship categories and controls for annual shifts in fleet composition and reporting

conditions. The label is descriptive and relative within the MRV dataset; it is not CII, EEXI, a certified energy-efficiency score, or a regulatory compliance label.




Figure 2. Distribution of ship-type-year tertile labels for CO₂ emissions per distance.

3.3. Feature Sets and Leakage Control

Two non-leakage feature sets are used for the main classification task. The `strict_static` set includes reporting year, ship type, parsed technical efficiency fields, port-of-registry information, home-port and ice-class indicators, and monitoring-method indicators. The `operational_no_emission` set adds non-emission operational variables such as time spent at sea and ice-navigation variables.

Target-family fields are excluded from classification. These include CO₂ intensity, total CO₂ emissions, fuel consumption, fuel intensity, transport-work intensity, and on-laden variants. The broader `consistency_screening` field set is reserved for the separate manual consistency-review task.

3.4. Temporal Evaluation and Models

The primary evaluation is temporal: models are trained on 2018–2021 labeled records, reported on a 2022 temporal holdout, tested on 2023 records, and checked externally on 2024 full-year labeled records. A random stratified 80/20 split over 2018–2023 is included only as an optimistic comparison. Practical use would require applying a model trained on earlier reports to later reports; therefore, temporal evaluation is the main evidence. The 2022 split is a reported temporal holdout, not a validation set used for hyperparameter tuning or model selection. All model configurations are fixed before evaluation. This distinction follows the broader concern that resampling designs for temporal, hierarchical, and repeated-observation data must match the intended generalization target [18,19].

Three classifiers are evaluated: Logistic Regression with class balancing, Random Forest with balanced subsampling, and Histogram-based Gradient Boosting. Random forests and gradient boosting are established tabular-learning baselines and have also been widely used in ship energy-consumption studies [12–14]. Numerical features are median-imputed. Categorical features are one-hot encoded for Logistic Regression and Random Forest and ordinal-encoded for Histogram-based Gradient Boosting. Macro-F1 and balanced accuracy are the primary metrics. The implementation uses Python and scikit-learn [20]. The core experimental configuration is summarized in Table 2.

Table 2. Core experimental configuration used for reproducible modeling.

Component	Setting	Notes
Random seed	42	Used for model training, random split, permutation importance, and residual cross-validation.
Logistic Regression	class_weight=balanced, solver=lbfgs, max_iter=800	Linear baseline with one-hot categorical encoding.
Random Forest	n_estimators=180, max_depth=16, min_samples_leaf=5	Uses balanced-subsample class weighting.
Histogram Gradient Boosting	learning_rate=0.05, max_iter=220, max_leaf_nodes=31	Uses ordinal categorical encoding and L2 regularization = 0.05.
Model selection	Fixed configurations	The 2022 holdout is reported, not used for hyperparameter search.
Random comparison	Stratified 80/20 split	Included only as an optimistic reference against temporal evaluation.
Consistency screening	Ship-type-specific models; minimum 300 eligible rows	Primary per-method threshold is 2%; sensitivity uses 1%, 2%, and 5%.
Residual screening	HGB regressor; 3-fold cross-validation	Predicts log-transformed CO ₂ per distance from consistency fields.

3.5. Ship-Type Ablation, Interpretability, and Consistency Screening

The ship-type ablation compares a unified model trained on all ship types with separate models trained for the five largest ship types: Bulk carrier, Oil tanker, Container ship, Chemical tanker, and General cargo ship. The ablation uses Histogram-based Gradient Boosting under both non-leakage feature sets.

Permutation importance is computed for the best temporal classifier using Macro-F1 drop after feature permutation. Error analysis focuses on medium-class records, because these records sit near tertile boundaries and are expected to have less stable class membership than the efficient and inefficient extremes.

The consistency-screening module is separate from classification. It uses complete annual reports with positive values for fuel consumption, fuel per distance, total CO₂ emissions, CO₂ per distance, and time spent at sea. Three methods are used: Isolation Forest, Local Outlier Factor, and regression residual ranking. The consensus screening score is the average of within-ship-type rank percentiles from the three methods. Candidate rows are described as consistency-review candidates for subsequent manual inspection. They are not interpreted as regulatory violations, fraud, false reporting, or externally confirmed findings. Sensitivity is assessed using 1%, 2%, and 5% per-method ship-type thresholds.

4. Results

4.1. Main Temporal Classification

The best temporal result is obtained by Histogram-based Gradient Boosting with the Operational feature set. On the 2023 test set, it reaches Macro-F1 = 0.603 and balanced accuracy = 0.605. The Static model is close, with Macro-F1 = 0.601 and balanced accuracy = 0.603, indicating that the added non-emission operational fields provide only a small improvement. The random stratified evaluation gives higher performance than the temporal evaluation. The best random split reaches Macro-F1 = 0.657, compared with Macro-F1 = 0.603 in the 2023 temporal test.

The temporal design does not remove all repeated-entity dependence. An IMO-overlap audit shows that 11,457 of 13,050 unique IMO numbers in the 2022 holdout also appear in the 2018–2021 training split, as do 10,439 of 12,615 in the 2023 test split and 10,029 of 12,905 in the 2024 external split (Table 3). Thus, the primary result is best interpreted as year-ahead repeated-fleet generalization. A stricter new-vessel generalization test would require grouped-by-IMO evaluation.

Table 3. IMO-number recurrence between the 2018–2021 training split and later-year evaluation splits.

Later Split	Later Unique IMO	Overlap IMO	Overlap Fraction
2022 holdout	13,050	11,457	0.878
2023 test	12,615	10,439	0.828
2024 external	12,905	10,029	0.777

Class-level metrics show where the temporal model is strongest and weakest. On the 2023 temporal test set, the best model reaches $F1 = 0.694$ for the efficient class, $F1 = 0.487$ for the medium class, and $F1 = 0.629$ for the inefficient class. The corresponding recalls are 0.745, 0.485, and 0.585, respectively. On the 2024 full-year external check, efficient-class $F1$ remains similar at 0.692, while medium and inefficient $F1$ decline to 0.479 and 0.584. This pattern confirms that the middle tertile is the least stable class and that the external-year decrease is concentrated mainly in the boundary and upper-intensity classes.

Table 4. Classification performance under temporal, external-year, and random stratified evaluation settings.

Split	Features	Model	Rows	Macro-F1	Balanced Accuracy
2023 test	Operational	HGB	12,615	0.603	0.605
2023 test	Static	HGB	12,615	0.601	0.603
2023 test	Operational	RF	12,615	0.560	0.573
2023 test	Static	RF	12,615	0.556	0.567
2024 full-year	Operational	HGB	12,905	0.585	0.587
2024 full-year	Static	HGB	12,905	0.580	0.582
Random test	Static	HGB	14,609	0.657	0.660
Random test	Operational	HGB	14,609	0.655	0.658

HGB = Histogram-based Gradient Boosting; RF = Random Forest; Static = `strict_static`; Operational = `operational_no_emission`.

Table 5. Class-level metrics for the best non-leakage temporal classifier.

Split	Class	Precision	Recall	F1	Support	Predicted Rows
2023 test	Efficient	0.649	0.745	0.694	4,210	4,833
2023 test	Medium	0.488	0.485	0.487	4,204	4,175
2023 test	Inefficient	0.681	0.585	0.629	4,201	3,607
2024 external	Efficient	0.647	0.743	0.692	4,314	4,953
2024 external	Medium	0.467	0.492	0.479	4,305	4,539
2024 external	Inefficient	0.659	0.525	0.584	4,286	3,413




Figure 3. Macro-F1 on the 2023 temporal test set for non-leakage feature sets and baseline classifiers.




Figure 4. Comparison between temporal evaluation and random stratified evaluation.


4.2. Ship-Type Ablation

Ship-type-specific modeling has heterogeneous effects. General cargo ship benefits under both feature sets, and Container ship improves under `strict_static`. Oil tanker shows smaller gains. Bulk carrier does not improve, and Container ship under `operational_no_emission` is slightly worse than the unified model.

Table 6. Unified versus ship-type-specific modeling for the five largest ship categories.

Ship Type	Features	Unified Macro-F1	Ship-Type Macro-F1	Delta
General cargo ship	Static	0.598	0.644	+0.046
Container ship	Static	0.609	0.646	+0.037
General cargo ship	Operational	0.593	0.626	+0.033
Chemical tanker	Operational	0.566	0.583	+0.017
Oil tanker	Static	0.671	0.687	+0.016
Oil tanker	Operational	0.674	0.682	+0.008
Bulk carrier	Operational	0.590	0.588	-0.003
Chemical tanker	Static	0.571	0.568	-0.003
Bulk carrier	Static	0.583	0.579	-0.004
Container ship	Operational	0.622	0.616	-0.006

Static = strict_static; Operational = operational_no_emission.


**Figure 5.** Macro-F1 comparison between unified and ship-type-specific models for the five largest ship types.

4.3. Feature Importance and Medium-Class Errors

Permutation importance identifies `technical_efficiency_value` as the strongest non-leakage feature, with a Macro-F1 drop of 0.228 after permutation. The next strongest predictors are `ship_type`, `technical_efficiency_type`, `port_of_registry`, and `time_spent_at_sea_hours`. The medium class is the main ambiguity source, which is consistent with the tertile-label design.

Table 7. Permutation importance and dominant medium-class error patterns.

Analysis	Rank	Item	Value
Permutation importance	1	Technical efficiency value	0.228
Permutation importance	2	Ship type	0.164
Permutation importance	3	Technical efficiency type	0.062
Permutation importance	4	Port of registry	0.024
Permutation importance	5	Time spent at sea	0.013
Medium error	1	Bulk carrier medium → efficient	339
Medium error	2	Container ship medium → inefficient	261
Medium error	3	General cargo ship medium → inefficient	176
Medium error	4	Bulk carrier medium → inefficient	169
Medium error	5	Container ship medium → efficient	161

**Figure 6.** Permutation importance of the best non-leakage temporal classifier.

4.4. Consistency-Review Screening and Sensitivity Analysis

The consistency-screening experiment models 85,932 complete annual rows across 14 ship types. At the 2% threshold, each method flags 1,727 rows. The all-three-method overlap contains 226 rows, while the at-least-two-method set contains 1,067 rows. These are manual consistency-review candidates, not externally confirmed findings or regulatory findings. The top 200 consensus candidates are mainly from Bulk carrier, Oil tanker, General cargo ship, Chemical tanker, and Container ship.

216

217

218

219


220

221

Table 8. Sensitivity of consistency-review candidates under 1%, 2%, and 5% ship-type-specific thresholds.

Thresh.	All Three	At Least Two	Any Method	Top 200 At Least Two	Top 200 All Three	Jaccard At Least Two	Jaccard All Three
1%	97	487	2,011	189	97	0.456	0.429
2%	226	1,067	3,888	200	182	1.000	1.000
5%	643	2,626	9,640	200	200	0.406	0.351

The public candidate table removes direct vessel identifiers and should be read as a pattern summary rather than a list of vessel-level findings. Within the top 200 consensus candidates, 182 rows are flagged by all three methods at the 2% threshold and 18 rows are flagged by two methods. The largest ship-type groups are Bulk carrier (59), Oil tanker (35), General cargo ship (22), Chemical tanker (19), and Container ship (17). Directional diagnostics split into low-intensity or low-volume patterns (77 rows), high-intensity or high-volume patterns (51 rows), and mixed patterns (72 rows). This diversity supports the use of a manual-review queue rather than a single deterministic anomaly rule.

**Figure 7.** Ship-type distribution of the top 200 consistency-review candidates.




Figure 8. Sensitivity of consistency-review overlap under 1%, 2%, and 5% thresholds.

5. Discussion

The results support three main findings. First, public MRV tables contain enough non-leakage information to support moderate temporal prediction of relative ship energy-efficiency strata. The performance level is not high enough to justify automated decision-making, but it is credible for screening, descriptive stratification, and decision-support analysis. The key methodological point is that performance is measured under temporal extrapolation rather than only under a random split.

Second, random stratified evaluation is optimistic for this dataset. A random split can mix similar reporting years and recurring ships across train and test partitions. The observed Macro-F1 gap between the random split and temporal test confirms that annual generalization should be the primary evaluation design for repeated regulatory-reporting datasets. However, the temporal split still includes substantial recurring-vessel overlap, so the present benchmark should not be overstated as new-vessel generalization.

Third, ship-type specialization is useful only for some categories. This is consistent with the competing effects of fleet heterogeneity and training-sample size. General cargo ship benefits from specialization, whereas Bulk carrier does not. The practical implication is that future applications should not assume a single modeling architecture is best for all ship categories.

The consistency-screening module should be interpreted cautiously. It uses variables that would be leakage-prone in classification, but those same variables are relevant for identifying records with unusual fuel, CO₂, distance, time, and intensity combinations. Because no external audit labels are available, the output is a prioritized manual-inspection list rather than proof of incorrect reporting. For this reason, public supplementary tables remove direct vessel identifiers and use candidate ranks and pattern fields instead of binding vessel identifiers to consistency-review judgments.

Several limitations remain. The tertile label is relative and data-internal; it is not equivalent to CIL, EEXI, or any compliance status. The data are annual aggregates and cannot capture voyage-level drivers such as speed, route, weather, loading, or port waiting time. The 2024 schema is broader than earlier schemas, so external-year interpretation requires care. The current temporal benchmark still contains recurring IMO numbers across years; future work should add grouped-by-IMO or new-vessel stress tests. Finally, the current models are deliberately lightweight; future work could compare specialized gradient-boosting libraries and could integrate AIS-derived operational variables where licensing permits.

6. Conclusions

This study presents MRV-EffScreen, a reproducible framework for public THETIS-MRV relative ship energy-efficiency stratification and emission consistency screening. Using 2018–2024 public records, the framework constructs ship-type-year relative labels, evaluates non-leakage features under temporal extrapolation, analyzes ship-type heterogeneity, and generates consensus consistency-review candidates. The best temporal classifier reaches Macro-F1 = 0.603 on the 2023 test set and Macro-F1 = 0.585 on the 2024 full-year external check, while random stratified evaluation gives a higher and more optimistic estimate. The consistency-screening module identifies stable manual-inspection candidates across threshold settings. The study demonstrates that public MRV data can support transparent machine-learning analysis when leakage control, temporal evaluation, and cautious interpretation are treated as first-order design requirements.

Author Contributions: Bo Li is the sole author and completed all CRediT roles: conceptualization, methodology, software, validation, formal analysis, investigation, resources, data curation, writing—original draft preparation, writing—review and editing, visualization, supervision, and project administration. The author has read and agreed to the published version of the manuscript.

Funding: This research received no external funding.

Institutional Review Board Statement: Not applicable.

Informed Consent Statement: Not applicable.

Data Availability Statement: The original public emission reports are available through the EU THETIS-MRV public emission report portal. The code, processing scripts, metadata, checksums, generated non-sensitive tables, figure files, and manuscript assets are available in the MRV-EffScreen repository at <https://github.com/TristanLib/MRV-EffScreen>. A Zenodo DOI will be added if the repository is archived after the first stable release.

Acknowledgments: The author thanks his family for their support. During the preparation of this manuscript and the associated codebase, AI-assisted tools were used to support drafting, code organization, language editing, and workflow documentation. The author reviewed and edited the outputs and takes full responsibility for the content of this publication.

Conflicts of Interest: The author declares no conflicts of interest.

Abbreviations

The following abbreviations are used in this manuscript:

AIS	Automatic Identification System
CII	Carbon Intensity Indicator
CO ₂	Carbon dioxide
EEXI	Energy Efficiency Existing Ship Index
EU MRV	European Union monitoring, reporting, and verification
IMO	International Maritime Organization
LOF	Local Outlier Factor
MRV	Monitoring, reporting, and verification

References

- European Maritime Safety Agency. THETIS-MRV. Available online: <https://emsa.europa.eu/thetis-mrv.html> (accessed on 20 May 2026).
- European Commission, Directorate-General for Climate Action. *2024 Report from the European Commission on CO₂ Emissions from Maritime Transport*; 2025. <https://doi.org/10.2834/8958061>.
- European Union. Regulation (EU) 2015/757 on the Monitoring, Reporting and Verification of Carbon Dioxide Emissions from Maritime Transport. 2015.
- European Union. Regulation (EU) 2023/957 Amending Regulation (EU) 2015/757 in Order to Provide for the Inclusion of Maritime Transport Activities in the EU Emissions Trading System and for the Monitoring,

- Reporting and Verification of Emissions of Additional Greenhouse Gases and Emissions from Additional Ship Types. 2023.
5. International Maritime Organization. *Fourth IMO Greenhouse Gas Study 2020*; 2020.
 6. International Maritime Organization. Revised GHG Reduction Strategy for Global Shipping Adopted. 2023. Available online: <https://www.imo.org/en/mediacentre/pressbriefings/pages/revised-ghg-reduction-strategy-for-global-shipping-adopted-.aspx> (accessed on 20 May 2026).
 7. International Maritime Organization. EEXI and CII: Ship Carbon Intensity and Rating System. Available online: <https://www.imo.org/en/mediacentre/hottopics/pages/eexi-cii-faq.aspx> (accessed on 20 May 2026).
 8. Yan, R.; Mo, H.; Wang, S.; Yang, D. Analysis and Prediction of Ship Energy Efficiency Based on the MRV System. *Marit. Policy Manag.* **2023**, *50*, 117–139. <https://doi.org/10.1080/03088839.2021.1968059>.
 9. Luo, X.; Yan, R.; Wang, S. After Five Years' Application of the European Union Monitoring, Reporting, and Verification (MRV) Mechanism: Review and Prospectives. *J. Clean. Prod.* **2024**, *434*, 140006. <https://doi.org/10.1016/j.jclepro.2023.140006>.
 10. Xing, H.; Chang, S.; Ma, R.; Wang, K. EU MRV Data-Based Review of the Ship Energy Efficiency Framework. *J. Mar. Sci. Eng.* **2025**, *13*, 1437. <https://doi.org/10.3390/jmse13081437>.
 11. Bullock, S.; Mason, J.; Broderick, J.; Larkin, A. Shipping and the Paris Climate Agreement: A Focus on Committed Emissions. *BMC Energy* **2020**, *2*, 5. <https://doi.org/10.1186/s42500-020-00015-2>.
 12. Fan, A.; Wang, Y.; Yang, L.; Tu, X.; Yang, J.; Shu, Y. Comprehensive Evaluation of Machine Learning Models for Predicting Ship Energy Consumption Based on Onboard Sensor Data. *Ocean Coast. Manag.* **2024**, *248*, 106946. <https://doi.org/10.1016/j.ocecoaman.2023.106946>.
 13. Breiman, L. Random Forests. *Mach. Learn.* **2001**, *45*, 5–32. <https://doi.org/10.1023/A:1010933404324>.
 14. Friedman, J.H. Greedy Function Approximation: A Gradient Boosting Machine. *Ann. Stat.* **2001**, *29*, 1189–1232. <https://doi.org/10.1214/aos/1013203451>.
 15. Liu, F.T.; Ting, K.M.; Zhou, Z.-H. Isolation Forest. In *Proceedings of the 2008 Eighth IEEE International Conference on Data Mining*; 2008; pp. 413–422. <https://doi.org/10.1109/ICDM.2008.17>.
 16. Breunig, M.M.; Kriegel, H.-P.; Ng, R.T.; Sander, J. LOF: Identifying Density-Based Local Outliers. In *Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data*; 2000; pp. 93–104. <https://doi.org/10.1145/335191.335388>.
 17. Kaufman, S.; Rosset, S.; Perlich, C. Leakage in Data Mining: Formulation, Detection, and Avoidance. *ACM Trans. Knowl. Discov. Data* **2012**, *6*, 1–21. <https://doi.org/10.1145/2382577.2382579>.
 18. Roberts, D.R.; Bahn, V.; Ciuti, S.; Boyce, M.S.; Elith, J.; Guillera-Arroita, G.; Hauenstein, S.; Lahoz-Monfort, J.J.; Schroeder, B.; Thuiller, W.; Warton, D.I.; Wintle, B.A.; Hartig, F.; Dormann, C.F. Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography* **2017**, *40*, 913–929. <https://doi.org/10.1111/ecog.02881>.
 19. Bergmeir, C.; Hyndman, R.J.; Koo, B. A note on the validity of cross-validation for evaluating autoregressive time series prediction. *Comput. Stat. Data Anal.* **2018**, *120*, 70–83. <https://doi.org/10.1016/j.csda.2017.11.003>.
 20. Pedregosa, F.; Varoquaux, G.; Gramfort, A.; Michel, V.; Thirion, B.; Grisel, O.; Blondel, M.; Prettenhofer, P.; Weiss, R.; Dubourg, V.; VanderPlas, J.; Passos, A.; Cournapeau, D.; Brucher, M.; Perrot, M.; Duchesnay, E. Scikit-learn: Machine Learning in Python. *J. Mach. Learn. Res.* **2011**, *12*, 2825–2830.